

Task 4: Calculate TF-IDF with MapReduce/Pig/Hive

1. Creating a table with the columns required to calculate TF-IDF and putting it into the local.
 - A. `hive> create table user_db.grouped_users_posts as select owneruserid as a, body as b, SUM(score) as c from user_db.stackdata_analysis group by owneruserid, body;`
 - B. `hive> INSERT OVERWRITE LOCAL DIRECTORY '/home/hduser/tfidfdata'
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
select owneruserid, body from user_db.stackdata_analysis where owneruserid in
(select id from user_db.grouped_users_posts);`
2. The MapReduce program needs the input file to be separated by space rather than comma. Therefore, replacing the comma with spaces and then moving the file into a folder in the local from HDFS.
 - A. `> cd tfidfdata`
 - B. `> sed 's/,/ /g' 000000_0 > inputfile`
 - C. `> hadoop fs -mkdir /mappred`
 - D. `> hadoop fs -put inputfile /mappred`
3. Now, the Implementation of TF-IDF in Hadoop using Python will be in three phases using three mappers and three reducers. Using the following commands to run TF-IDF on the cluster.
 - A. `> hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar -file
/home/hduser/MAPPER_01.py /home/hduser/REDUCER_01.py -mapper "python
MAPPER_01.py" -reducer "python REDUCER_01.py" -input hdfs://cluster-3299-
m/mapinput/inputfile -output hdfs://cluster-3299-m/mappred1`
 - B. `> hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar -file
/home/hduser/MAPPER_02.py /home/hduser/REDUCER_02.py -mapper "python
MAPPER_02.py" -reducer "python REDUCER_02.py" -input hdfs://cluster-3299-
m/mappred1/part-00000 hdfs://cluster-3299-m/mappred1/part-00001
hdfs://cluster-3299-m/mappred1/part-00002 hdfs://cluster-3299-
m/mappred1/part-00003 hdfs://cluster-3299-m/mappred1/part-00004 -
output hdfs://cluster-3299-m/mappred2`

```
C. > hadoop jar /usr/lib/hadoop-mapreduce/hadoop-streaming.jar -file
/home/hduser/MAPPER_03.py /home/hduser/REDUCER_03.py -mapper "python
MAPPER_03.py" -reducer "python REDUCER_03.py" -input hdfs://cluster-3299-
m/mappred2/part-00000 hdfs://cluster-3299-m/mappred2/part-00001
hdfs://cluster-3299-m/mappred2/part-00002 hdfs://cluster-3299-
m/mappred2/part-00003 hdfs://cluster-3299-m/mappred2/part-00004 -
output hdfs://cluster-3299-m/mappredf
```

- Using the `hadoop fs -getmerge` command to merge the output files into a single CSV file (`tfidfout.csv`). Now, replacing the spaces and saving the file into another CSV file(`tfidfout1.csv`).

```
> sed -e 's/\s/,/g' tfidfout.csv > tfidfout1.csv
```

5. Thereafter, creating an external table in HIVE and loading the CSV file into the table.

```
A. hive> create external table if not exists TFIDF_data2 (Term String, Id int, tfidf float)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ',';
```

B. `hive> load data local inpath 'tfidfout1.csv' overwrite into table TFIDF_data2;`

6. Finally, running the following query to calculate the per-user TF-IDF of the top 10 terms for each of the top 10 users.

```
hive> SELECT *
FROM (
SELECT ROW_NUMBER()
OVER(PARTITION BY Id
ORDER BY tfidf DESC) AS TfidfRank, *
FROM TfIDF_data2) n
WHERE TfidfRank IN (1,2,3,4,5,6,7,8,9,10);
```

[illegible]